

Car Price Prediction with Machine Learning

Objective

The goal of this project is to develop a machine learning model that predicts car prices based on user input. Students will work with a real-world [dataset](#) from Kaggle, exploring data, building predictive models, and deploying an API.

Project Steps

1. Dataset Exploration

- Go to Kaggle to learn more about the dataset via the link below:
<https://www.kaggle.com/datasets/sidharth178/car-prices-dataset>
- Load the dataset into a Jupyter Notebook.
- Perform exploratory data analysis: check missing values, data distribution, correlations, and outliers.
- Clean and preprocess the data.
- Split the data into train, validation and test dataset.

2. Baseline Model

- Define a simple baseline model.
- Choose an appropriate performance metric.
- Evaluate the baseline model.

3. Model Comparison

- Experiment with multiple machine learning models.
- Compare their performance using the chosen metric.
- Interpret the results.

4. Hyperparameter Tuning

- Use GridSearchCV to optimize model hyperparameters.
- Get the best hyperparameters.

5. Model Saving

- Select the best-performing model with the package pickle.

6. API Development with FastAPI

- Create an API that allows users to input car features and output a predicted price.
- Implement at least one endpoint for prediction.
- Test the API locally with sample requests.

7. Code Versioning with GitHub

- Organize the code structure and control the environment.
- Use Git for version control and push the project to GitHub.
- Write a README file explaining the project and its usage.

8. Deployment on Render

- Deploy the FastAPI application on Render.
- Ensure the API is accessible online.

Deliverables

- A Jupyter Notebook with data analysis, baseline model, model comparisons, and hyperparameter tuning.

- A FastAPI script for the prediction API in a GitHub repository with all project files.
- A deployed API link on Render.
- A short report (maximum 5 pages) explaining the workflow and results.

In the report you have to present:

- The context and the project,
- Why you use the metric of performance,
- The performance of your algorithm,
- The best hyperparameter,
- The explanation of how work your API,
- The link of the different resources of the project (notebook, github, url of the API on render).
- Conclusion of the project.

All deliverables must be submitted by March 3 at my email address
gautherotmorgan0@gmail.com

This project will serve as the final evaluation for the course, testing both technical skills and the ability to apply data science to a real-world problem. 🚀